



EDB

Postgres® for the AI Generation

Postgres on Kubernetes Workshop

Requirements for running the hands-on locally

Have these things ready prior to joining the workshop

- Bash shell available (i.e. either running Linux or Mac, or WSL on Windows laptops)
- Internet access (for downloading code and images)
- Have at least a little knowledge of Kubernetes and Postgres (youtube is a great resource)
- Docker is installed
- Run without errors
 - `$ docker ps`
- k3d or kind is installed (k3d preferred as it seems there are less issues with it)
- Create local Kubernetes cluster (run either) - remove it again (we don't need it)
 - `$ kind create cluster --name cnpq-workshop`
 - `$ k3d cluster create cnpq-workshop`



Agenda

| Session |
|---|
| Introduction to Kubernetes |
| Introduction to Operators and CNPG |
| DBA Mindset Shift - how Kubernetes changes the game (and doesn't) |
| Interactive session |
| Additional functionality in CNPG |
| A shameless plug + Discussion |
| Lunch |



Introduction to Kubernetes



Cloud Native!!

Let's play Buzzword Bingo!

Benefits:

Increased Agility, Improved Scalability, Higher Resilience, Optimization of resources, Lower Time to Market, ...

Core terms:

Decoupling, Configuration as Code, Declarative APIs, Immutable architecture, Services Meshes, Change Management (DevOps procedures)

Implementation details:

Microservices, Containers, Orchestration, Cloud Agnostic, CI/CD, CoDe, ... (DevOps???)



From Servers to Containers - why Kubernetes matters

A new way to manage workloads

- From physical servers (bare metal) to virtual servers (VMs)
 - Full operating system
- Containers
 - Shares the OS Kernel with the Host system but runs in complete isolation.

A container is essentially a package that contain everything an application needs to run.

(Think of it as a small, self-contained box that holds application code and dependencies)

Run in isolation using Linux Namespaces and use cgroups to put cap resources.

- Docker - the beginning of it all
 - Which is why the 'code' used to create a container is often called a "Dockerfile"
 - The official name today though is "Containerfile"



Example of a Containerfile

This is what creates a 'Container'

```
# Use the official PostgreSQL image as a base.  
FROM postgres:16.2-alpine  
  
# Set environment variables for the database configuration.  
ENV POSTGRES_USER=admin  
ENV POSTGRES_PASSWORD=securepassword  
ENV POSTGRES_DB=mydb  
  
# Copy our custom initialization script and our custom startup script into the container.  
# init-schema.sql: Contains your database schema (e.g., CREATE TABLE statements).  
# start.sh: Our custom bash script to run the database.  
COPY ./init-schema.sql /docker-entrypoint-initdb.d/  
COPY ./start.sh /usr/local/bin/  
  
# Make the start.sh script executable.  
RUN chmod +x /usr/local/bin/start.sh  
  
# The CMD instruction now tells the container to execute our custom bash script.  
CMD ["start.sh"]
```



How Kubernetes uses a container

- How do you run a container on Kubernetes?
 - You create a Pod with the container inside

A Pod is the smallest unit in Kubernetes. It can contain a number of containers.

The containers are either the actual application or something that initialises the application.
Or, they can add functionality, like network filtering, log collection, ...



The magic in Kubernetes (Operator or Control Loop)

- Once you have a pod, you can deploy that to Kubernetes
- Either as:

Single Pod - meaning no special handling

Deployment - a way to automatically scale to multiple pods and apply HA

If you tell Kubernetes to run a deployment with 3 replicas of your Pod it will do so, and **maintain that count even if one of the pods fail**. Pods are automatically killed and created if Kubernetes detects a problem with a pod.

This is what is known as a Kubernetes Control Loop - also known as an “Operator”.

Kubernetes natively has one Operator - but we can add our own (ones that understand our app)



Why is Kubernetes special?

- Declarative deployments vs. imperative
 - Kubernetes tends to be declarative (though some things can be done imperatively - you really shouldn't)
- Leans into Infrastructure as Code, DevOps etc.
- Perfect for MicroServices etc.

But why run the databases in Kubernetes?

- Database workload benefit from Kubernetes as well (automatic healing, scale in/out, etc.)
- One platform to run them all - no “context switch” between different environments



Introduction to CloudNativePG and EDB



What is CNPG?

CNPG = Cloud Native Postgres (<https://cloudnative-pg.io/>)

CloudNativePG is the Kubernetes operator that covers the full lifecycle of a highly available PostgreSQL database cluster with a primary/standby architecture, using native streaming replication.

- An Operator (Control Loop) plus the images needed to run databases using the operator

CNPG is designed bottom-up for Kubernetes. This sets it apart from other available operators where the architecture is Patroni-based and then patched to ‘work’ in Kubernetes.



Quick facts about CNPG

- Immutable images (secure by default)
 - Only user available when opening a shell: postgres or enterpriseDb (depending on database type)
- Heavily tested and vulnerability tested by EDB prior to release
- CNPG uses StorageClass (if you want)
- You don't manage StatefulSets or Persistent Volumes or even Persistent Volume Claims
 - CNPG does that for you
- Maturity level for CNPG is the highest possible: **AutoPilot**

But wait ... which storage provider should we use?

- Databases need fast storage. Until someone builds a Network Storage Solution faster than local storage ...
You should use local storage - but you can configure something else



Interactive session

It's time to go hands-on!



Use case

The environment



Features shown during the demo

- Installation operator and CNPG plugin
- Check the CloudNativePG operator status
- Installation Backup (Barman) Operator
- Setup PostgreSQL Cluster on Kubernetes
- Create the backup of the PostgreSQL Cluster
- Run out-of-the-place recovery
- Run switchover
- Simulate failover szenario
- Run minor upgrade using rolling upgrade approach
- Scale out/down
- Test fencing
- Test hibernation
- Database migration
- Major upgrade (in-place)

Deployment

High Availability

Administration

Monitoring

Backup and Recovery

Patching

Fencing

Hibernation

Last CloudNativePG tested version is 1.26



Monitoring

This demo is in



<https://github.com/EnterpriseDB/cnpg-hands-on>

Download or pull the code

cd to the new directory.



Setup the demo env



Call to action: Create k3d cluster

- In the terminal 1:
 - Go to the directory /home/workshop/workshop/cnp-demo:
 - Create the k3d cluster - run the script:
`./00_start_infra.sh`
 - Check the cluster:
`kubectl get nodes`
`kubectl get pods`



Call to action: Run Minio server

- In the terminal 1:
 - Go to the directory /home/workshop/workshop/cnp-demo:
 - Start MinIO:
`./start_minio_docker_server.sh &`



Call to action: Connect to the MinIO server

- In the browser open the new tab and go to
 - <http://localhost:9001>
 - Connect as user **admin** with the password: **password**
- The page will appear:

The screenshot shows the MinIO Object Store web interface. On the left, there is a dark sidebar with the MinIO logo at the top. Below it, under the heading 'User', are links for 'Object Browser' (which is highlighted in blue), 'Access Keys', and 'Documentation'. Under the heading 'Administrator', are links for 'Buckets' (which is also highlighted in blue) and 'Policies'. The main content area has a light background. At the top, it says 'Object Browser'. On the right side of the header are three small icons: a question mark, a moon, and a refresh symbol. Below the header, there is a large rectangular box containing a section titled 'Buckets'. The title 'Buckets' is preceded by a icon of three horizontal bars. The text inside the box explains: 'MinIO uses buckets to organize objects. A bucket is similar to a folder or directory in a filesystem, where each bucket can hold an arbitrary number of objects.' At the bottom of this box, there is a link 'To get started, [Create a Bucket.](#)'.

Use case Plug-in installation



The “cnpg” plugin for kubectl

- The official CLI for CloudNativePG
 - Available also as RPM or Deb package
- Extends the ‘kubectl’ command:
 - Customize the installation of the operator
 - Status of a cluster
 - Perform a manual switchover (promote a standby) or a restart of a node
 - Issue TLS certificates for client authentication
 - Declare start and stop of a Kubernetes node maintenance
 - Destroy a cluster and all its PVC
 - Fence a cluster or a set of the instances
 - Hibernate a cluster
 - Generate jobs for benchmarking via pgbench and fio
 - Issue a new backup
 - Start pgadmin



Name: cluster-example
Namespace: default
System ID: 7100921006673293335
PostgreSQL Image: ghcr.io/cloudnative-pg/postgresql:14.3
Primary instance: cluster-example-2
Status: Cluster in healthy state
Instances: 3
Ready instances: 3
Current Write LSN: 0/C000060 (Timeline: 4 - WAL File: 000000040000000000000000C)

Certificates Status

| Certificate Name | Expiration Date | Days Left Until Expiration |
|-----------------------------|-------------------------------|----------------------------|
| cluster-example-replication | 2022-08-21 13:15:00 +0000 UTC | 89.95 |
| cluster-example-server | 2022-08-21 13:15:00 +0000 UTC | 89.95 |
| cluster-example-ca | 2022-08-21 13:15:00 +0000 UTC | 89.95 |

Continuous Backup status

First Point of Recoverability: 2022-05-23T13:37:08Z
Working WAL archiving: OK
WALs waiting to be archived: 0
Last Archived WAL: 00000004000000000000000B @ 2022-05-23T13:42:09.37537Z
Last Failed WAL: -

Streaming Replication status

| Name | Sent LSN | Write LSN | Flush LSN | Replay LSN | Write Lag | Flush Lag | Replay Lag | State | Sync State | Sync Priority |
|-------------------|-----------|-----------|-----------|------------|-----------|-----------|------------|-----------|------------|---------------|
| cluster-example-3 | 0/C000060 | 0/C000060 | 0/C000060 | 0/C000060 | 00:00:00 | 00:00:00 | 00:00:00 | streaming | async | 0 |
| cluster-example-1 | 0/C000060 | 0/C000060 | 0/C000060 | 0/C000060 | 00:00:00 | 00:00:00 | 00:00:00 | streaming | async | 0 |

Instances status

| Name | Database Size | Current LSN | Replication role | Status | QoS | Manager Version |
|-------------------|---------------|-------------|------------------|--------|------------|-----------------|
| cluster-example-3 | 33 MB | 0/C000060 | Standby (async) | OK | BestEffort | 1.15.0 |
| cluster-example-2 | 33 MB | 0/C000060 | Primary | OK | BestEffort | 1.15.0 |
| cluster-example-1 | 33 MB | 0/C000060 | Standby (async) | OK | BestEffort | 1.15.0 |



Call to action: Install CNPG plugin

- In the terminal 1:
 - Install CNPG plugin:

```
./01_install_plugin.sh
```



Use case Operator installation



Call to action: Operator Installation

- Install the operator
- Check the installed CNPG Operator in the terminal
- In the terminal 1:
 - Install CNPG operator:

`./02_install_operator.sh`

- Check the installation of the operator (try several times):

`./03_check_operator_installed.sh`



Use case Backup Configuration



Backup and Recovery - Backup Methods

- CNPG Operator supports several backup methods:
 - Database backup to the object storage - with Barman tool
 - Until the CNPG Version 1.26 barman was integrated into CNPG Operator
 - From version 1.26 onwards, EDB has provided a comprehensive backup operator based on Barman technology - **Barman Cloud Plugin**
 - Native volume snapshots
- WAL management
 - Object store
- Physical Base backups
 - Object store
 - Kubernetes level backup integration (Velero/OADP, Veem Kasten K10, generic interface)*
 - Kubernetes Volume Snapshots



Backup and Recovery - Choosing the best ssolution

| Feature | Object Store | Volume Snapshots |
|-------------------------------|--------------|----------------------|
| WAL archiving | Required | Recommended |
| Cold backup | ✗ | ✓ |
| Hot backup | ✓ | ✓ |
| Incremental copy | ✗ | ✓ |
| Differential copy | ✗ | ✓ |
| Backup from a standby | ✓ | ✓ |
| Snapshot recovery | ✗ | ✓ |
| Retention policies | ✓ | ✗ |
| Point-in-Time Recovery (PITR) | ✓ | Requires WAL archive |
| Underlying technology | Barman Cloud | Kubernetes API |



Backup and Recovery - Barman Cloud Plugin

- Features:
 - Continuous physical backup on “backup object stores”
 - Scheduled and on-demand base backups
 - Continuous WAL archiving (including parallel)
 - Backup can be created from primary or a standby
 - Backups encryption
 - Support for parallel backups
 - Support for backup compression
 - Support for recovery window retention policies (e.g. 30 days)



Backup and Recovery - Barman Cloud Plugin

- Recovery means creating a new cluster starting from a “recovery object store”
 - Then pull WAL files (including in parallel) and replay them
 - Full (End of the WAL) or PITR
- Both rely on Barman Cloud technology
 - AWS S3
 - Azure Storage compatible
 - Google Cloud Storage
 - MinIO



Backup and Recovery - Kubernetes Volume Snapshot

- Transparent support for:
 - Incremental backup and recovery at block level
 - Differential backup and recovery at block level
 - Based on copy on write
- Leverage the storage class to manage the snapshots, including:
 - Data mobility across network (availability zones, Kubernetes clusters, regions)
 - Relay files on a secondary location in a different region, or any subsequent one
 - Encryption
- Enhances Very Large Databases (VLDB) adoption



Backup & Recovery via Snapshots: some numbers

Let's now talk about some initial benchmarks I have performed on volume snapshots using 3 `r5.4xlarge` nodes on AWS EKS with the `gp3` storage class. I have defined 4 different database size categories (tiny, small, medium, and large), as follows:

| Cluster name | Database size | pgbench init scale | PGDATA volume size | WAL volume size | pgbench init duration |
|---------------|---------------|--------------------|--------------------|-----------------|-----------------------|
| <i>tiny</i> | 4.5 GB | 300 | 8 GB | 1 GB | 67s |
| <i>small</i> | 44 GB | 3,000 | 80 GB | 10 GB | 10m 50s |
| <i>medium</i> | 438 GB | 3,0000 | 800 GB | 100 GB | 3h 15m 34s |
| <i>large</i> | 4,381 GB | 300,000 | 8,000 GB | 200 GB | 32h 47m 47s |

The table below shows the results of both backup and recovery for each of them.

| | Cluster name | 1st backup duration | 2nd backup duration after 1hr of pgbench | Full recovery time |
|--|---------------|---------------------|--|--------------------|
| | <i>tiny</i> | 2m 43s | | 4m 16s |
| | <i>small</i> | 20m 38s | | 16m 45s |
| | <i>medium</i> | 2h 42m | | 2h 34m |
| | <i>large</i> | 3h 54m 6s | | 2m 2s |

<https://www.enterprisedb.com/postgresql-disaster-recovery-with-kubernetes-volume-snapshots-using-cloudnativepg>



Call to action: Barman Cloud Operator configuration

- Installing cert-manager
- Installing Barman Cloud operator
- Configuration the storage in MinIO
- In the terminal 1:
 - Install Barman Operator:
`./04_install_barman_plugin.sh`
 - Check the installation of the barman operator:
`./05_check_barman_plugin.sh`



Use case

Create the postgres cluster



Synchronizing the state of a Postgres database

- Being a DBMS, PostgreSQL is a stateful workload in Kubernetes
- Stateless workloads achieve HA and DR mainly through traffic redirection
- Stateful workloads require the state to be replicated in multiple locations:
 - **Storage-level** replication
 - **Application-level** replication (in our case, application = Postgres)
- Postgres has a very robust and powerful native replication system
 - We've built it
 - Founded on the Write Ahead Log
 - Read-only standby servers
 - Supports also synchronous replication controlled at the transaction level
- **We recommend application-level** over storage-level replication for Postgres



Bootstrap - different ways of creating a cluster

- Create a new cluster from scratch
 - “initdb”: named after the standard “initdb” process in PostgreSQL that initializes an instance
- Create a new cluster from an existing one:
 - Directly (“pg_basebackup”), using physical streaming replication
 - Directly (logical backup/restore) using pg_dump and pg_restore
 - Indirectly (“recovery”), from an object store
 - To the end of the WAL
 - Can be used to start independent replica clusters in continuous recovery
 - Using PITR



Storage management

- Storage is the most critical component for a database
- Direct support for Persistent Volume Claims (PVC)
 - We deliberately do not use Statefulsets
- The PVC storing the PGDATA is central to CloudNativePG
 - Our motto is: “PGDATA is worth a 1000 pods”
- Storage agnostic
- Freedom of choice
 - Local storage
 - Network storage
- Automated generation of PVC
- Support for PVC templates
 - Storage classes



Call to action: Configure and Install the Postgres cluster

- In the terminal 1:

- Create the yaml file:

```
./06_get_cluster_config_file.sh
```

- Create the postgres cluster:

```
./07_install_cluster.sh
```

- In the terminal **2** - Check the Postgres cluster status::

```
./08_show_status.sh
```



Call to action: Create table test with 1000 rows

- Once cluster is running ... (minimum the primary) insert the data - run the script:

```
./09_insert_data.sh
```



Use case Backup & Restore



Backup demonstration

- Create the full backup
- Check Backup in MinIO UI
- Restore the database from the backup



Call to action: Connect to the MinIO server

- In the browser open the new tab and go to
 - `http://<vm ip>:9001`
 - Connect as user **admin** with the password: **password**
- The page will appear:

The screenshot shows the MinIO Object Store interface. On the left, there's a sidebar with navigation links: User (Object Browser, Access Keys, Documentation), Administrator (Buckets, Policies), and a bottom section with a 'Logout' link. The main area is titled 'Object Browser' and shows a bucket named 'cnp'. The bucket was created on 'Mon, Apr 14 2025 18:54:07 (GMT+2)' and has 'PRIVATE' access. It contains 63.6 MiB of data across 6 objects. A search bar at the top right says 'Start typing to filter objects in the bucket'. Below the bucket details, there's a table with two rows. The first row is collapsed (indicated by a minus sign) and the second row is expanded, showing a folder named 'cluster-example'. The table has columns for 'Name' (sorted by Last Modified) and 'Last Modified'.

| Name | Last Modified |
|-----------------|---------------|
| cluster-example | |

- Click on cluster-example and check the backup of WAL files



Call to action: Create the full backup

- With this step we will:
 - Create the full backup of the postgres cluster in the MinIO storage:
- In the terminal 1:
 - Run the script:

```
./10_backup_cluster.sh
```

- Check the backup status:

```
./11_backup_describe.sh
```

- Check the created backup in the MinIO GUI



Call to action: Check Backup in MinIO UI

- Check the created backup in the MinIO GUI

The screenshot shows the MinIO Object Browser interface. On the left, there's a sidebar with navigation links: User (Object Browser, Access Keys, Documentation), Administrator (Buckets, Policies, Identity), and a license notice for AGPLv3. The main area is titled 'Object Browser' and shows a bucket named 'cnp'. The bucket was created on Monday, April 14, 2025, at 18:54:07 (GMT+2). It has PRIVATE access and contains 63.6 MiB of data across 6 objects. A search bar at the top right says 'Start typing to filter objects in the bucket'. Below the bucket details, there's a table with two rows: one for the bucket itself and one for a folder named 'cluster-example'.

| Name | Last Modified |
|-----------------|---------------|
| cnp | |
| cluster-example | |



Call to action: Restore the database from the backup

- With this step we will:
 - Create the new cluster cluster-restore
 - Restore the full backup created in the previous step in the new cluster:
- In the terminal 1:
 - Run the restore:
`./12_restore_cluster.sh`
 - Check the creation status:
`kubectl get pods -w # after creation stop the execution with <ctrl>+c`
 - Check the table test in the cluster-restore, run the script:
`./13_check_restore.sh`
 - Delete the cluster-restore to avoid resource problems during the workshop:
`kubectl delete cluster cluster-restore`



Use case Promote & Failover



Call to action: Promote standby to the primary

- With this step we will:
 - promote a pod in the cluster to primary, so you can start with maintenance work or test a switch-over situation in your cluster
 - Promote is the option of the cnpg plugin
- In the terminal 1:
 - Run the script:
`./14_promote.sh`
- In the terminal **2**:
 - Check the failover cluster status:
`./08_show_status.sh`



Call to action: Run failover test

- With this step we will:
 - Delete the primary database of the cluster cluster-example
 - The operator will:
 - detect failure of the primary
 - promote standby to the new primary and create the new standby according to the configuration (3 instances)
- In the terminal 1:
 - Run the script:
`./15_failover.sh`
- In the terminal **2**:
 - Check the failover cluster status:
`./08_show_status.sh`



Use case Minor Upgrade



Rolling updates

- Update of a deployment with ~zero downtime
 - Standby servers are updated first
 - Then the primary:
 - supervised / unsupervised
 - switchover / restart
- When they are triggered:
 - Security update of Postgres images
 - Minor update of PostgreSQL
 - Configuration changes when restart is required
 - Update of the operator
 - Unless in-place upgrade is enabled



Call to action: Run the Promote and Upgrade

- With this step we will:
 - Run the postgres minor update from the version 16.4 to 16.5
- In the web terminal **2** (prepare a terminal for status - and one to run the admin-commands):
 - Check the upgrade status:
`./08_show_status.sh`
- In the terminal **1**:
 - Run the script:
`./16_minor_upgrade.sh`
- In terminal **2**: (prepare a terminal for status - and one to run the admin-commands):
 - check Postgres version: PostgreSQL Image: [quay.io/enterprisedb/postgresql:16.5](https://quay.io/repository/enterprisedb/postgresql)



Use case Scale-out and scale-down



Scale up and down of replicas

- The operator allows you to scale up and down the number of instances in a PostgreSQL cluster.
- New replicas are started up from the primary server and participate in the cluster's HA infrastructure.
- The CRD declares a "scale" subresource that allows you to use the `kubectl scale` command.



Call to action: Scale-out the postgres cluster

- With this step we will:
 - Add the 1 standby to the cluster
- In the web terminal 1:
 - Run the script:
`./18_scale_out.sh` (using `-replicas=X...` another way would be to update the YAML)
- In the web terminal **2**:
 - Check the cluster status:
`./08_show_status.sh`



Call to action: Scale-down the postgres cluster

- With this step we will:
 - Remove 2 standby pods from the cluster
- In the web terminal 1:
 - Run the script:
`./19_scale_down.sh`
- In the web terminal **2**:
 - Check the cluster status:
`./08_show_status.sh`



Use Case Fencing



Fencing

- **Fencing** is the process of protecting the data in one, more, or even all instances of a PostgreSQL cluster when they appear to be malfunctioning.
- When an instance is fenced, the PostgreSQL server process is guaranteed to be shut down, while the pod is kept running.
- This ensures that, until the fence is lifted, data on the pod isn't modified by PostgreSQL and that you can investigate file system for debugging and troubleshooting purposes.



Call to action: Stop postgres process on the pod

- In the web terminal 1:

- Run the script:

- ```
./20_fencing.sh on
```

- In the web terminal **2**:

- Check the cluster status:

- ```
./08_show_status.sh
```



Call to action: Start the postgres process on the pod

- In the terminal 1:

- Run the script:

```
./20_fencing.sh off
```

- In the terminal **2**:

- Check the cluster status:

```
./08_show_status.sh
```



Use case Hibernation



Hibernation

- CloudNativePG supports **hibernation** of a running PostgreSQL cluster in a declarative manner
 - through the cnpg.io/hibernation annotation
- Hibernation enables saving CPU power by removing the database pods while keeping the database PVCs
- This feature simulates scaling to 0 instances.



Call to action: Stop the postgres cluster

- In the terminal 1:
 - Run the script:
`./21_hibernation.sh on`

- In the terminal **2**:
 - Check the cluster status:
`./08_show_status.sh`



Call to action: Start the postgres cluster

- In the terminal 1:
 - Run the script:
`./21_hibernation.sh off`
- In the terminal **2**:
 - Check the cluster status:
`./08_show_status.sh`



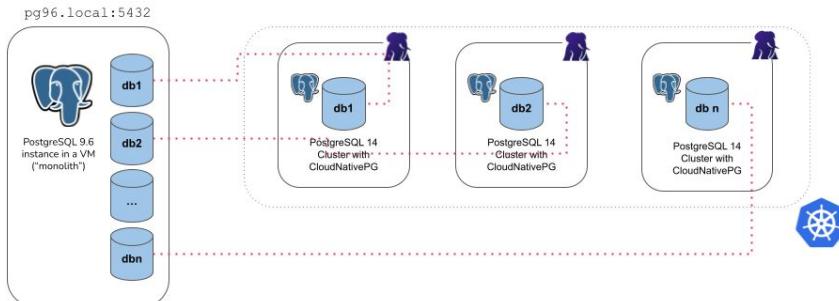
Use case Database Migration



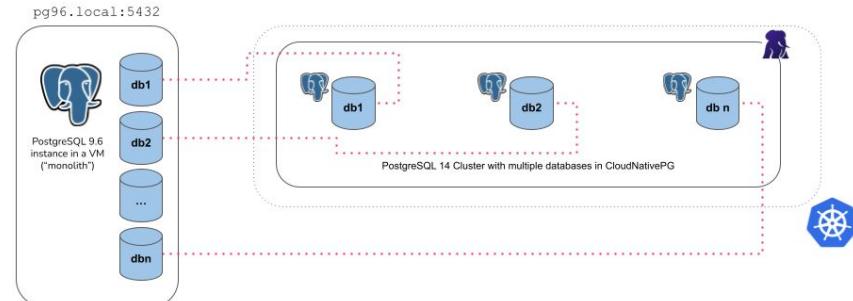
Database Migration

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “import” in the bootstrap section
- The operator uses internally postgres tools pg_dump and pg_restore
- This method can be used to **migrate** another database or to run **out-of-the-place upgrade**
- Possible settings:

Microservices:



Monolith:



Call to action: Migrate the cluster example to the new cluster

- In the terminal 1:
 - Create the new cluster v17 and migrate the database app from the cluster cluster-example::
`./22_major_upgrade_by_copy.sh`
 - Connect to the migrated cluster and check version and data::
`./23_verify_data_migrated_16_17.sh`



Use case Major Upgrade



Database Major Upgrade

- CNPG supports 3 types of major upgrade:
 - **In-place**: you should test carefully this method before you migrate your production database. You should create the fallback scenario in case of failure
 - **Copy**: we discussed this topic in “Database migration”
 - **Logical replication**: you can create the new database in the new cluster and replicate the data between two clusters.
 - CloudNativePG enhances this capability by providing declarative support for key PostgreSQL logical replication objects:
 - Publications via the Publication resource
 - Subscriptions via the Subscription resource



Call to action: In-place major upgrade

- In this step we will:
 - upgrade the existing cluster cluster-example from the Postgres version 16 to 17
 - choose in-place method

Note: For this action the downtime should be planned: during the upgrade process postgres instances are not available!

- In the terminal 1:
 - Run in-place major upgrade for the cluster cluster-example:
`./24_major_upgrade_in_place.sh`
 - Connect to the upgraded cluster and check version:
`./25_verify_major_upgrade_16_17.sh`



Use case Monitoring



Call to action: Setup monitoring

- In the terminal 1:
 - Change the directory::
 - Install the prometheus rules:
`./monitoring/01_prometheus_rules.sh`
 - Start port forwarding for prometheus and grafana:
`./monitoring/02_port_forwarding_prometheus_grafana.sh`
- Download the Grafana Dashboard to your laptop:
 - <https://github.com/cloudnative-pg/grafana-dashboards/blob/main/charts/cluster/grafana-dashboard.json>



Call to action: Explore Prometheus

- In the browser open the new tab and go to
 - <http://<vm ip>:9090>
- The prometheus page will appear - search for “cnpq” metrics:

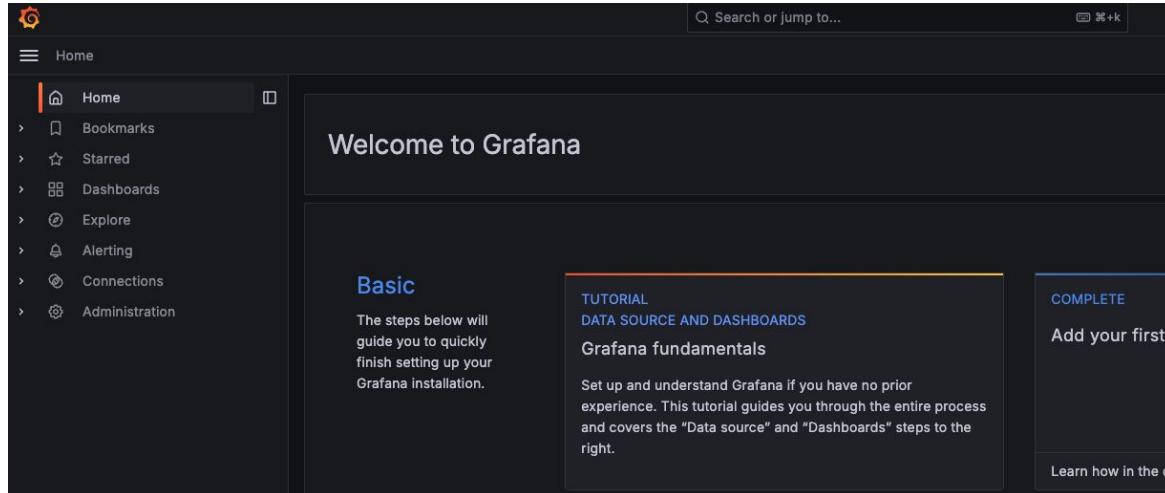
The screenshot shows the Prometheus web interface. At the top, there's a navigation bar with icons for Prometheus, Query, Alerts, and Status. Below the navigation bar is a search bar containing the text "cnpq". A green arrow points to this search bar. To the right of the search bar, a list of metrics is displayed. Each metric entry consists of a blue icon, a metric name, a type indicator (e.g., gauge, counter), and a brief description. The metrics listed include:

| Metric Name | Type | Description |
|--|---------|--|
| cnpq_backends_max_tx_duration_seconds | gauge | Maximum duration of a transaction in seconds |
| cnpq_backends_total | gauge | |
| cnpq_backends_waiting_total | gauge | |
| cnpq_collector_collection_duration_seconds | gauge | |
| cnpq_collector_collections_total | counter | |
| cnpq_collector_fencing_on | gauge | |
| cnpq_collector_first_recoverability_point | gauge | |
| cnpq_collector_last_available_backup_timestamp | gauge | |
| cnpq_collector_last_collection_error | gauge | |
| cnpq_collector_last_failed_backup_timestamp | gauge | |
| cnpq_collector_lo_pages | gauge | |
| cnpq_collector_manual_switchover_required | gauge | |
| cnpq_collector_nodes_used | gauge | |
| cnpq_collector_pg_wal | gauge | |
| cnpq_collector_pg_wal_archive_status | gauge | |
| cnpq_collector_postgres_version | gauge | |
| cnpq_collector_replica_mode | gauge | |



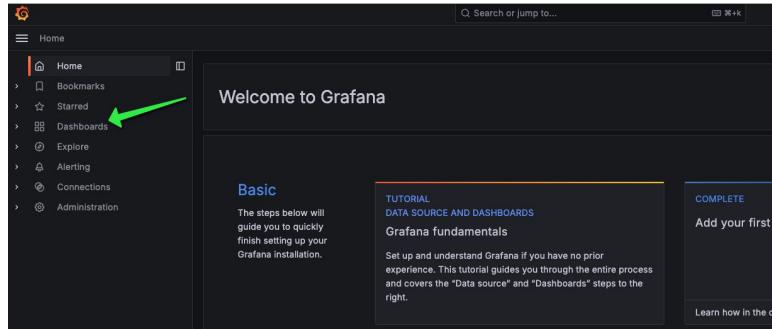
Call to action: Access the Grafana page

- In the browser open the new tab and go to
 - <http://<vm ip>:3000>
 - Connect as user **admin** with the password: **prom-operator**
- The grafana page will appear

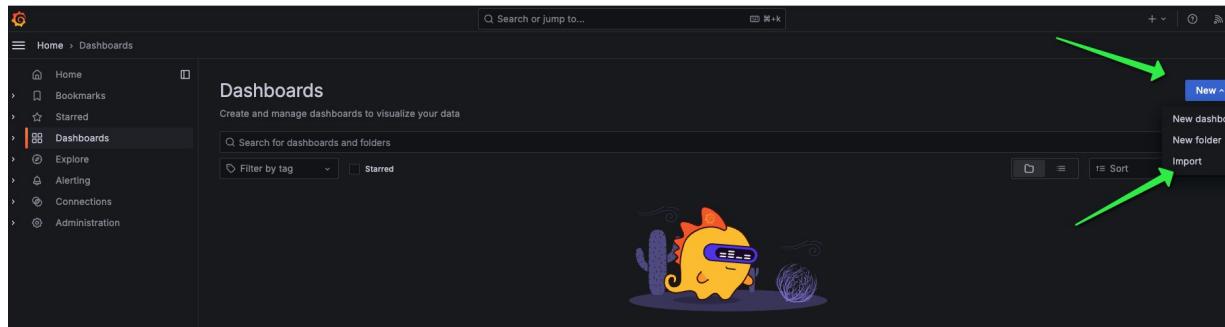


Call to action: Configure Grafana

- Go to Dashboards

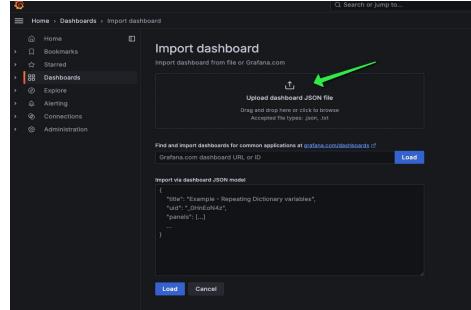


- Press "New", then "Import":

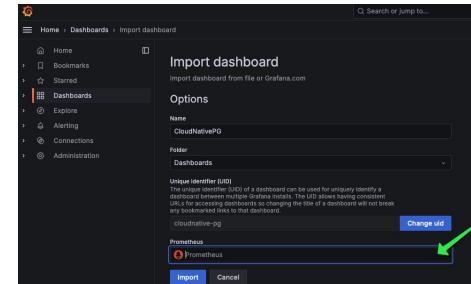


Call to action: Configure Grafana - continued

- Upload the Dashboard json file:



- Upload grafana-dashboard.json file
- Select Prometheus as the data source:



Call to action: Explore CNPG Dashboard- continued

- Explore the CNPG Dashboard:



What more?
(some additional features from EDB)



What we didn't show you today

- PgBouncer (Pooler) integration
 - Create a PgBouncer deployment and automatically configure to the cluster.
- TLS
 - CNPG supports TLS/SSL for encrypted connections etc.
 - The cluster you're running right now uses certificates for streaming replication
- RBAC
 - The operator uses a dedicated service account 'cnpq-manager' to interact with K8s
 - Of course, within the database, everything is as 'normal'
- Auditing
 - PGAudit can be configured by changing the parameters
.spec.postgresql.parameters.pgaudit.*
 - The operators automatically load PGAudit and runs "CREATE EXTENSION"



Advanced Security



Password policy management

DBA managed password profiles, compatible with Oracle profiles



Audit compliance

Track and analyze database activities and user connections



Virtual private databases

Fine grained access control limits user views



EDB/SQL protect

SQL firewall, screens queries for common attack profiles



Data redaction

Protect sensitive information for GDPR, PCI and HIPAA compliance



Code protection

Protects sensitive IP, algorithms or financial policies

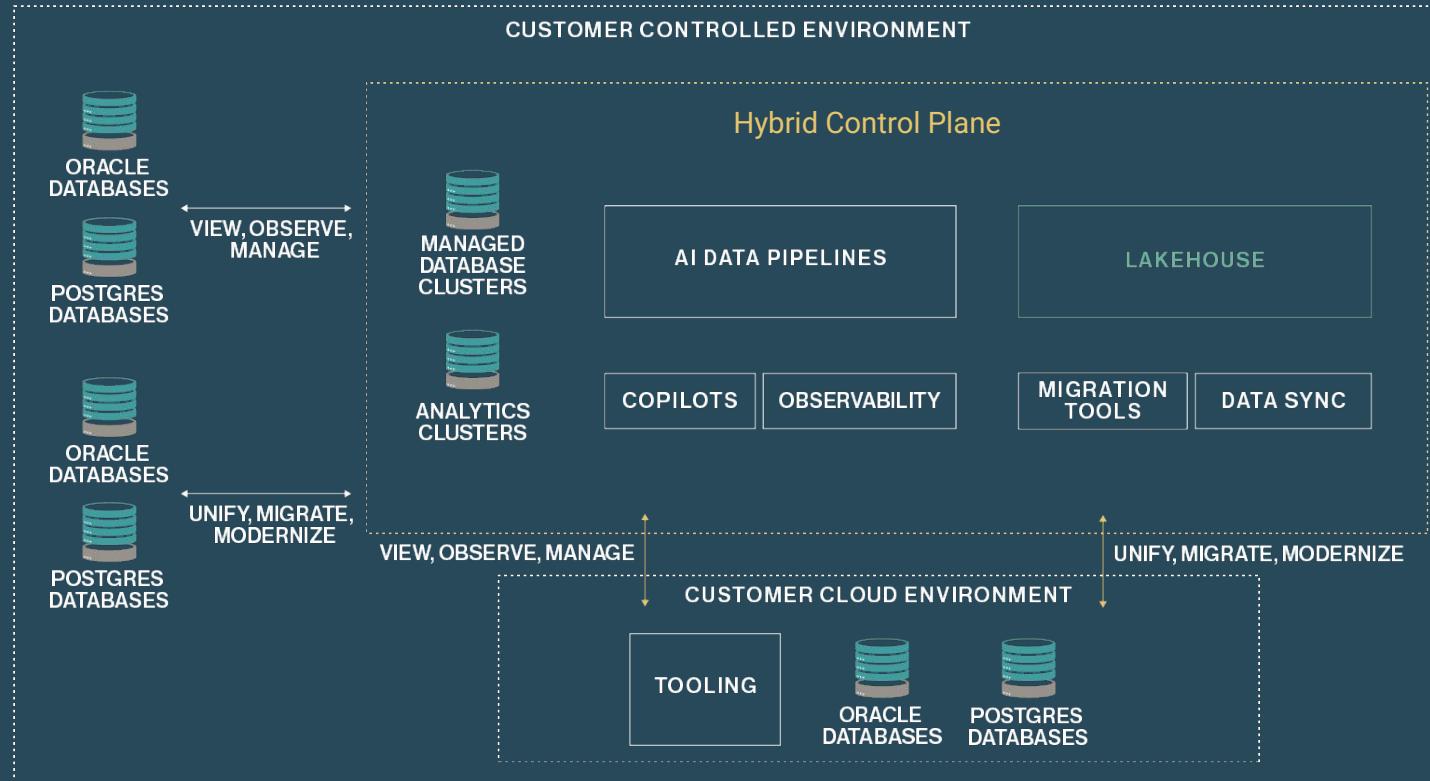


Transparent Data Encryption (EDB-only features)

- Transparent Data Encryption (TDE) is a feature of EDB Postgres Advanced Server and EDB Postgres Extended Server that prevents unauthorized viewing of data in operating system files on the database server and on backup storage
- Data encryption and decryption is managed by the database and does not require application changes or updated client drivers
- EDB Postgres Advanced Server and EDB Postgres Extended Server provide hooks to key management that is external to the database allowing for simple passphrase encrypt/decrypt or integration with enterprise key management solutions, with initial support for:
 - Amazon AWS Key Management Service (KMS)
 - Google Cloud - Cloud Kay Management Service
 - Microsoft Azure Key Vault
 - HashiCorp Vault (KMIP Secrets Engine and Transit Secrets Engine)
 - Thales CipherTrust Manager
- Data will be unintelligible for unauthorized users if stolen or misplaced



Hybrid Control Plane at a glance



EDB Postgres® AI (PG:AI)

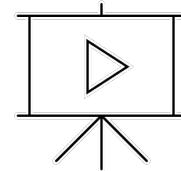
What if you could do all this, in a single install?

- **Database Observability**
 - Monitor all your Postgres databases using internal API, Prometheus and Grafana
- **Database migrations**
 - Migrate Oracle database directly into the platform
- **Database Provisioning**
 - Run Postgres databases as simple as point and click
- **Analytics Accelerator**
 - 30x performance on analytics queries compared to Postgres
 - Query data through Postgres (include both live-data, analytics and Delta Table content)
 - Automatically manage live data -> analytics storage -> delta table
- **AI Accelerator**
 - Manage and deploy AI models
 - Build low-code AI bots (and more advanced ones, of course)
 - Automatic data management of files and databases for AI Knowledge bases



EDB Postgres® AI

LIVE DEMO



BACKUP SLIDES



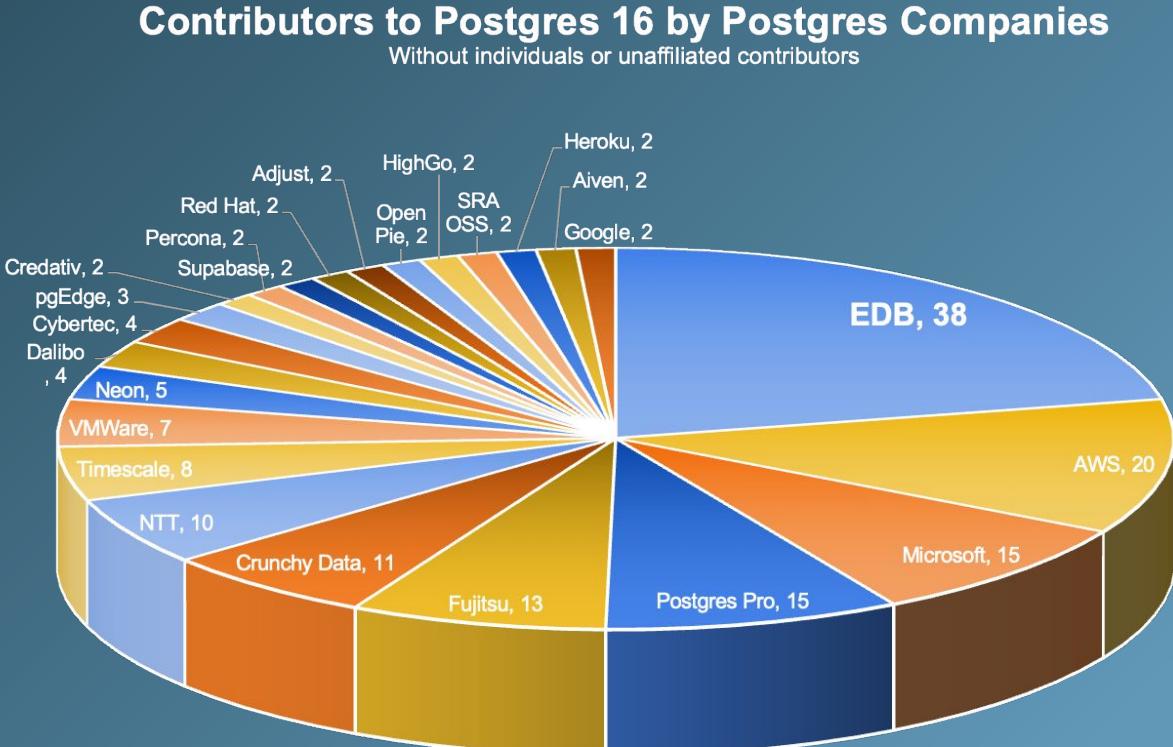
20+ years of Postgres innovation & adoption

- Number one contributor to Postgres, fastest-growing and most loved Database in the world
 - 2 Core Team members, 7 Committers, 9 Major Contributors, 10 Contributors, #1 site for desktop downloads
- Over 700 employees in more than 30 countries
- EDB Postgres AI
 - The industry's first platform that can be deployed as cloud, software or physical appliance
 - Secure, compliant and enterprise grade performance guaranteed



Large Developer Community

- 407 + contributors to Postgres 16
- 111 companies actively contributing to Postgres 16



Postgres has won the database race



2023:
45.5%



2024:
48.7%

Stack Overflow Survey 2023/2024



BANKING FINANCIAL

TECHNOLOGY

TELCO



BBVA

AON



**AMERICAN
EXPRESS**



Santander



IAG
Insurance
Australia
Group



SONY

tomtom

Braintree
a PayPal service



RSA



SAMSUNG

Postmates



NOKIA

DELL EMC

AT&T



OPTUS

T Systems

VONAGE



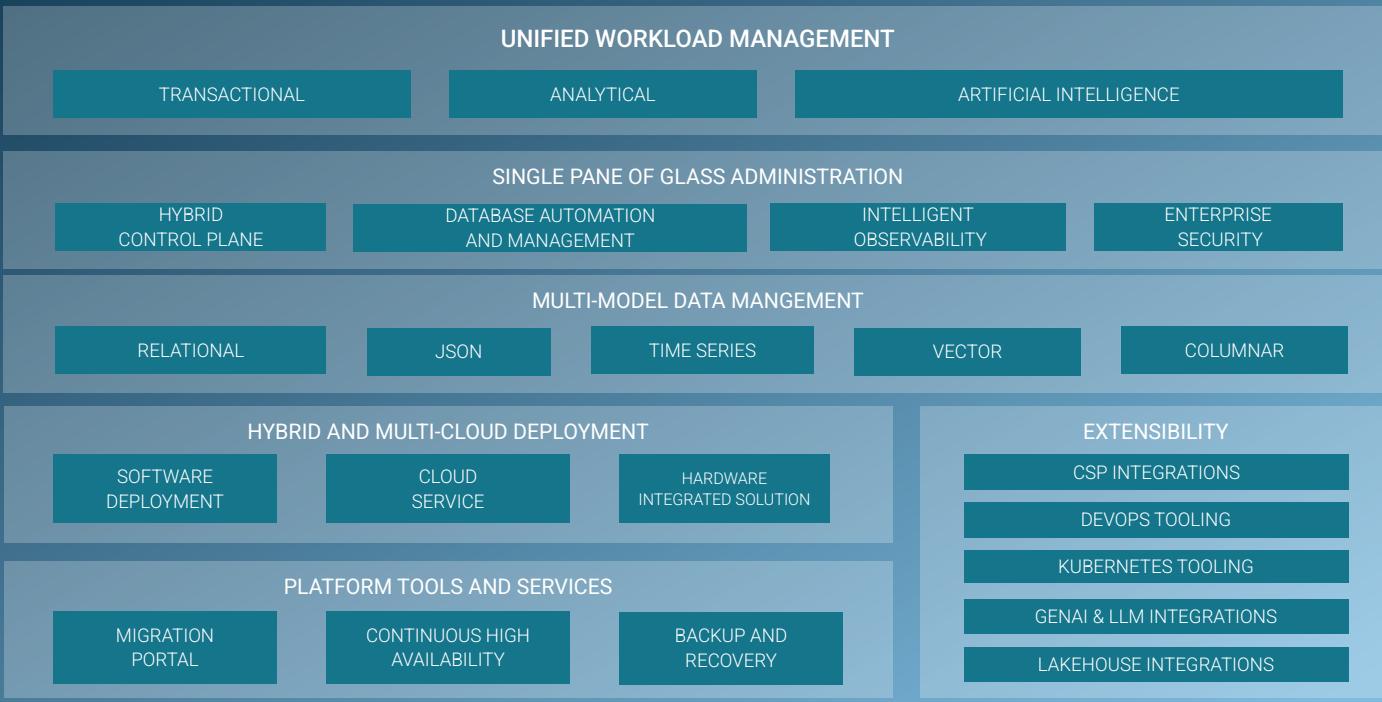
Telefónica

vodafone

Telstra

verizon

EDB POSTGRES AI PLATFORM



CNPG Operator: Reference Architecture and functionalities

Kubernetes timeline

- 2014, June: Google open sources Kubernetes
- 2015, July: Version 1.0 is released
- 2015, July: Google and Linux Foundation start the CNCF
- 2016, November: The operator pattern is introduced in a blog post
- 2018, August: The Community takes the lead
- 2019, April: Version 1.14 introduces **Local Persistent Volumes**
- 2019, August: EDB team starts the Kubernetes initiative
- 2020, June: we publish this blog about benchmarking local PVs on bare metal
- 2020, June: Data on Kubernetes Community founded
- 2021, February: EDB Cloud Native Postgres (CNP) 1.0 released
- 2022, May: **EDB donates CNP** and open sources it under CloudNativePG
- 2025, January: CloudNativePG was recognized as an official **#CNCF** project



Enabling the same PostgreSQL everywhere

From self-managed to fully managed DBaaS in the Cloud

- Same applications
- Faster innovation
- Performance and scalability
- Stability, security and control
- Seamless integration
- Obsolescence



Private



Hybrid



Multi-cloud



Public



Bare Metal



Virtual Machines



Containers



A kubernetes operator for Postgres



Kubernetes adoption is rising and it is already the de facto **standard** orchestration tool



PostgreSQL clusters “management the kubernetes way” enables many cloud native usage patterns, e.g. spinning up, disposable clusters during tests, one cluster per microservice and one database per cluster

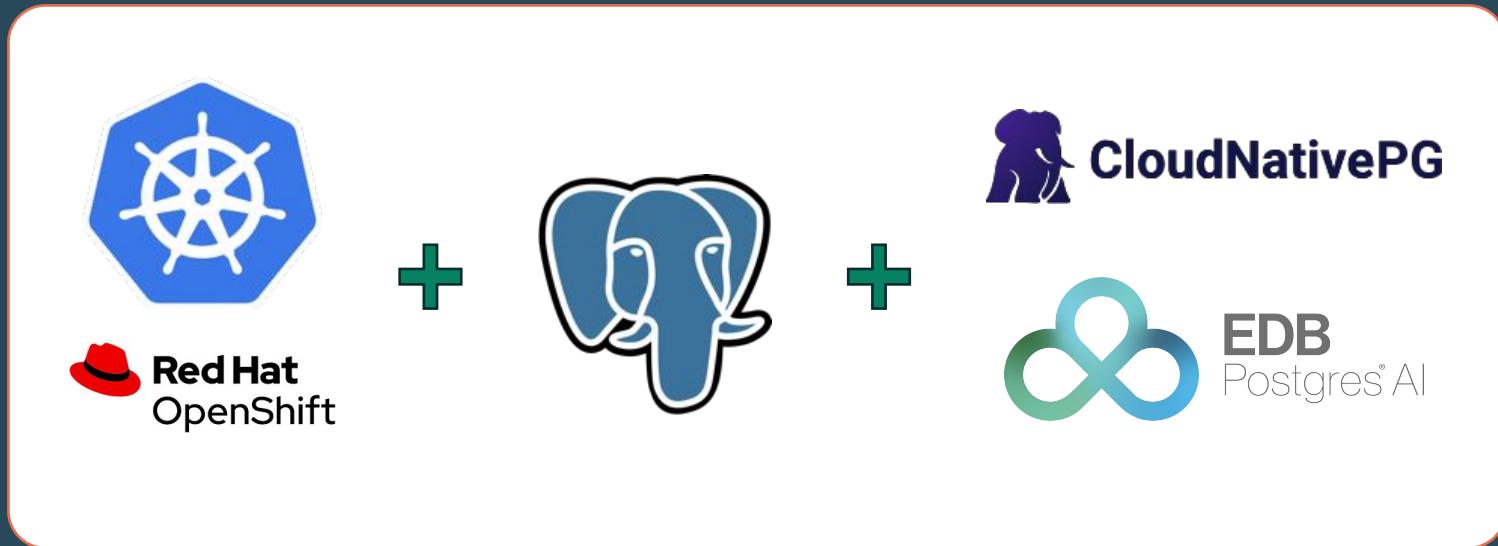


CNPG tries to encode years of experience managing PostgreSQL clusters into **an Operator which should automate all the known tasks** a user could be willing to do

Our PostgreSQL operator must simulate the work of a DBA



Win Technology



Autopilot

It automates the steps that a human operator would do to deploy and to manage a Postgres database inside Kubernetes, including automated failover.



Security

A grayscale photograph of a security guard from behind. The guard is wearing a light-colored zip-up jacket with the word "SECURITY" printed in large, bold, white capital letters on the back. He is holding a dark walkie-talkie up to his ear with his right hand. The background is a plain, light-colored wall.

CloudNativePG is secured by default.





It doesn't rely on statefulsets and uses its own way to manage persistent volume claims where the PGDATA is stored.

Data persistence



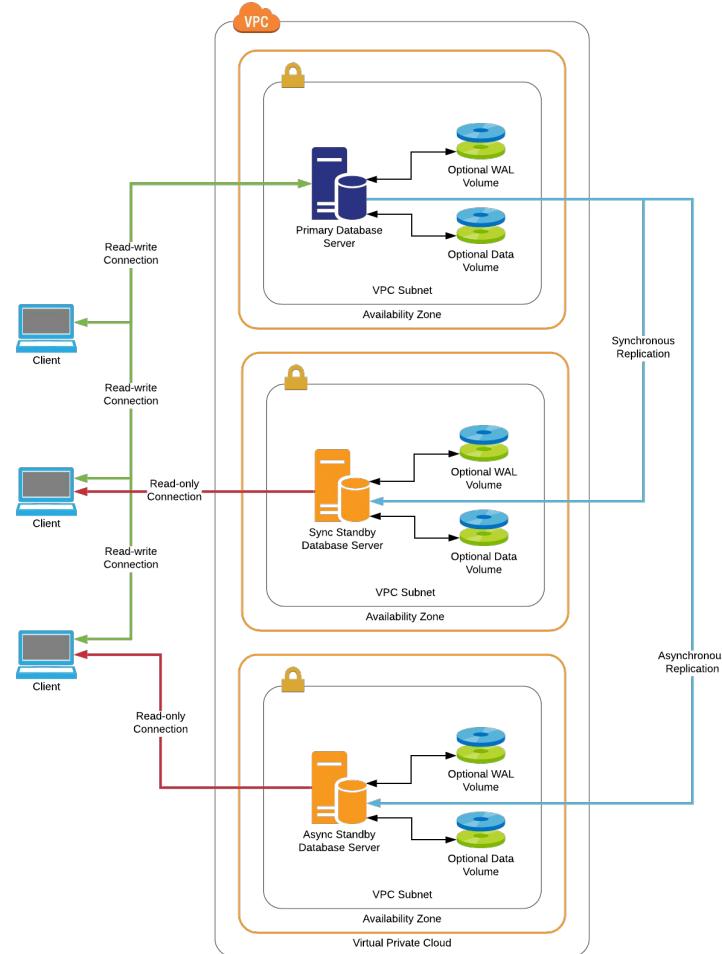
Designed for Kubernetes

It's entirely declarative, and directly integrates with the Kubernetes API server to update the state of the cluster — for this reason, it does not require an external failover management tool.



Imperative vs Declarative

- Create and configure VMs
- Create a PostgreSQL 13 instance
- Configure for replication
- Clone a second one
- Set it as a replica
- Clone a third one
- Set it as a replica
- Configure networking
- Configure security
- etc.



Convention over configuration

Declarative - simple to install, simple to maintain

There's a PostgreSQL 17 cluster with 2 replicas:

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: myapp-db
spec:
  instances: 3
  imageName: quay.io/enterprisedb/postgresql:17

storage:
  size: 10Gi
```



Features

| Deployment | Administration | Backup & Recovery | Monitoring | Security | High Availability |
|---------------------|--------------------------|-------------------|-----------------------------|---------------------|------------------------|
| Kubernetes operator | Single node | Backup | Prometheus | TDE | Switchover |
| Kubernetes plugin | Cluster (Multi node) | Recovery | Grafana dashboards | Certificates | Failover |
| EDB Postgres (EPAS) | PostgreSQL configuration | PITR | Postgres Enterprise Manager | Data redaction | Scale out / scale down |
| PostGIS | Pooling | Volume Snapshots | Logging | Password management | Minor / Major updates |



Decision-making for choosing the deployment platform



Advantage of deploying Postgres Databases in Kubernetes

Automation & Orchestration

- 01 |
- Self-healing
 - Automated scaling
 - Rolling updates

Self-healing

- 02 |
- Best resource utilization
 - Dynamic Resource allocation

Rolling updates

- 03 |
- Cloud-agnostic
 - Consistent deployment

Service discovery & networking

- 04 |
- Built-it service discovery
 - Load Balancing

Automated backups and disaster recovery

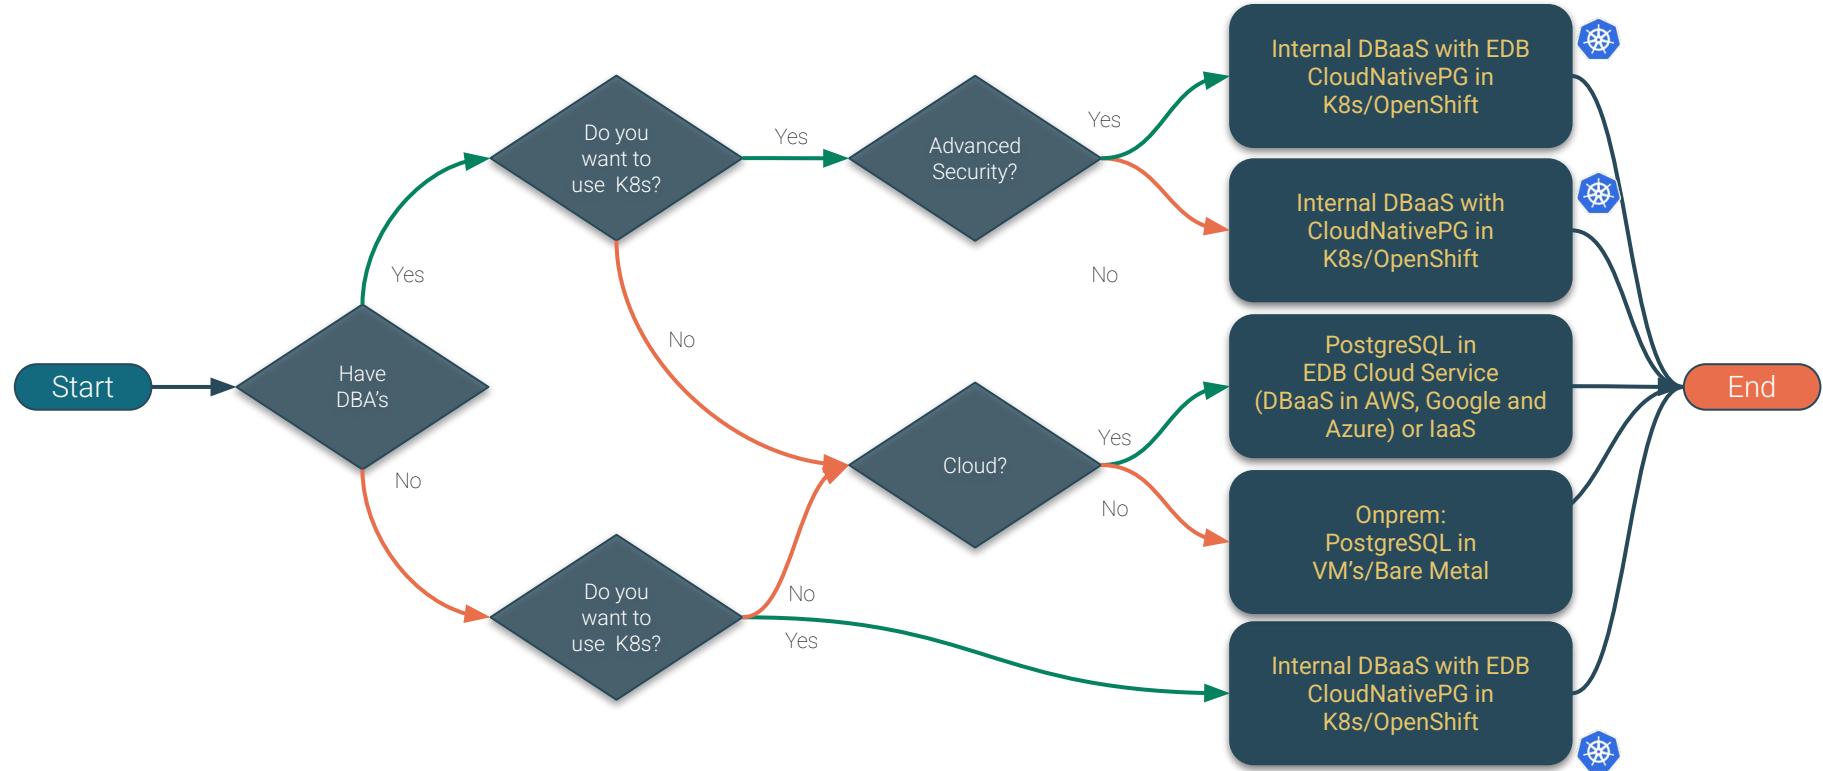
- 05 |
- Automated backups
 - Multi-region failover

Security & access control

- 06 |
- RBAC
 - Secret management



Decision-making for choosing the deployment platform



Architectures



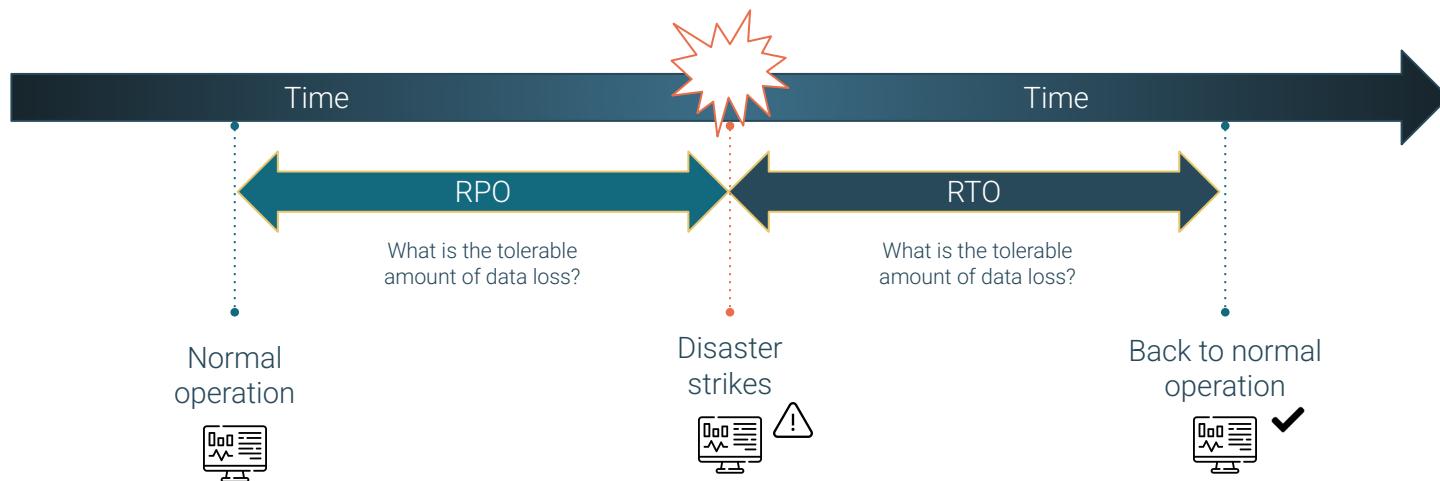
When to choose Kubernetes over VMs?

- 01 |** Cloud Native Applications that already run in Kubernetes
- 02 |** Scalable, replicated databases
- 03 |** Applications requiring automated failover and self-healing
- 04 |** Teams skilled in Kubernetes who want a unified infrastructure



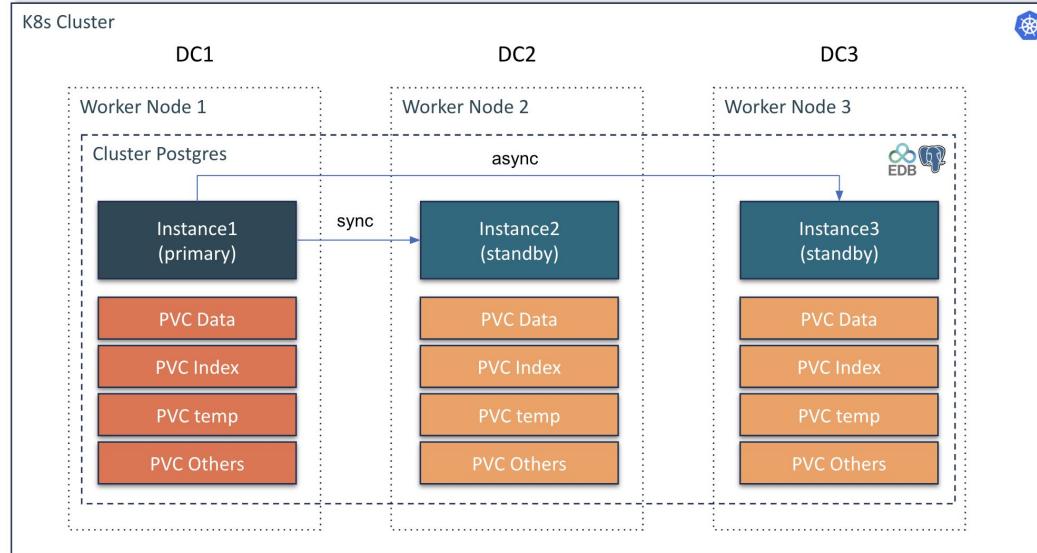
Concepts

- Recovery Point Objective (**RPO**) and Recovery Time Objective (**RTO**) are key concepts in disaster recovery and business continuity planning, particularly related to data loss and system downtime.



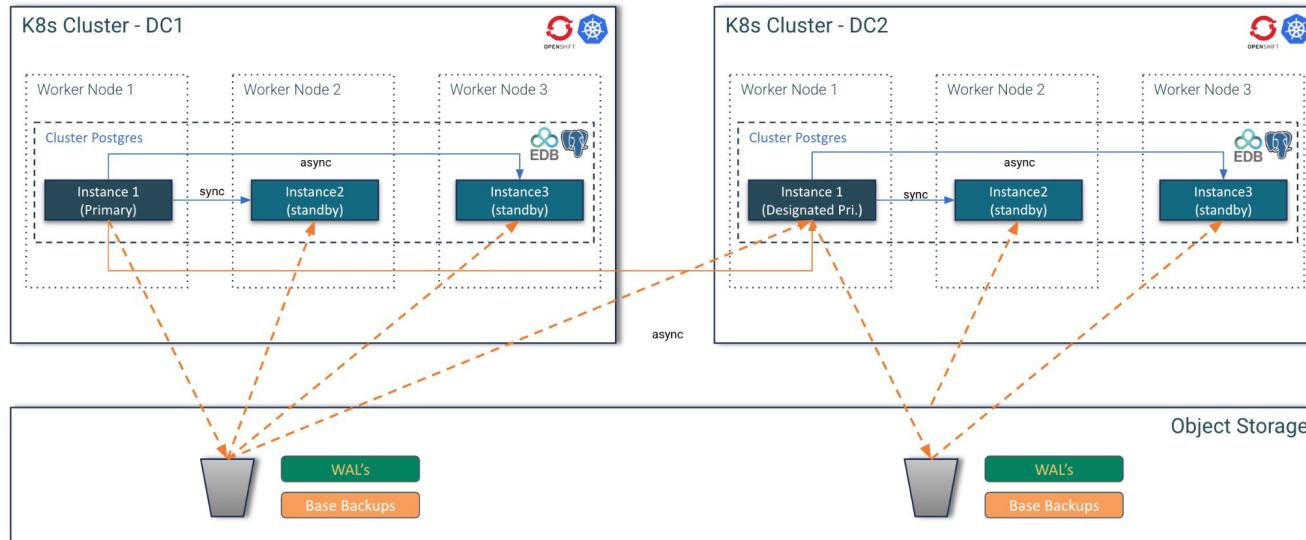
Red Hat Recommendation

Red Hat recommend stretched clusters ONLY when latencies don't exceed 5 milliseconds (ms) round-trip time (RTT) between the nodes in different locations, with a maximum RTT of 10 ms.

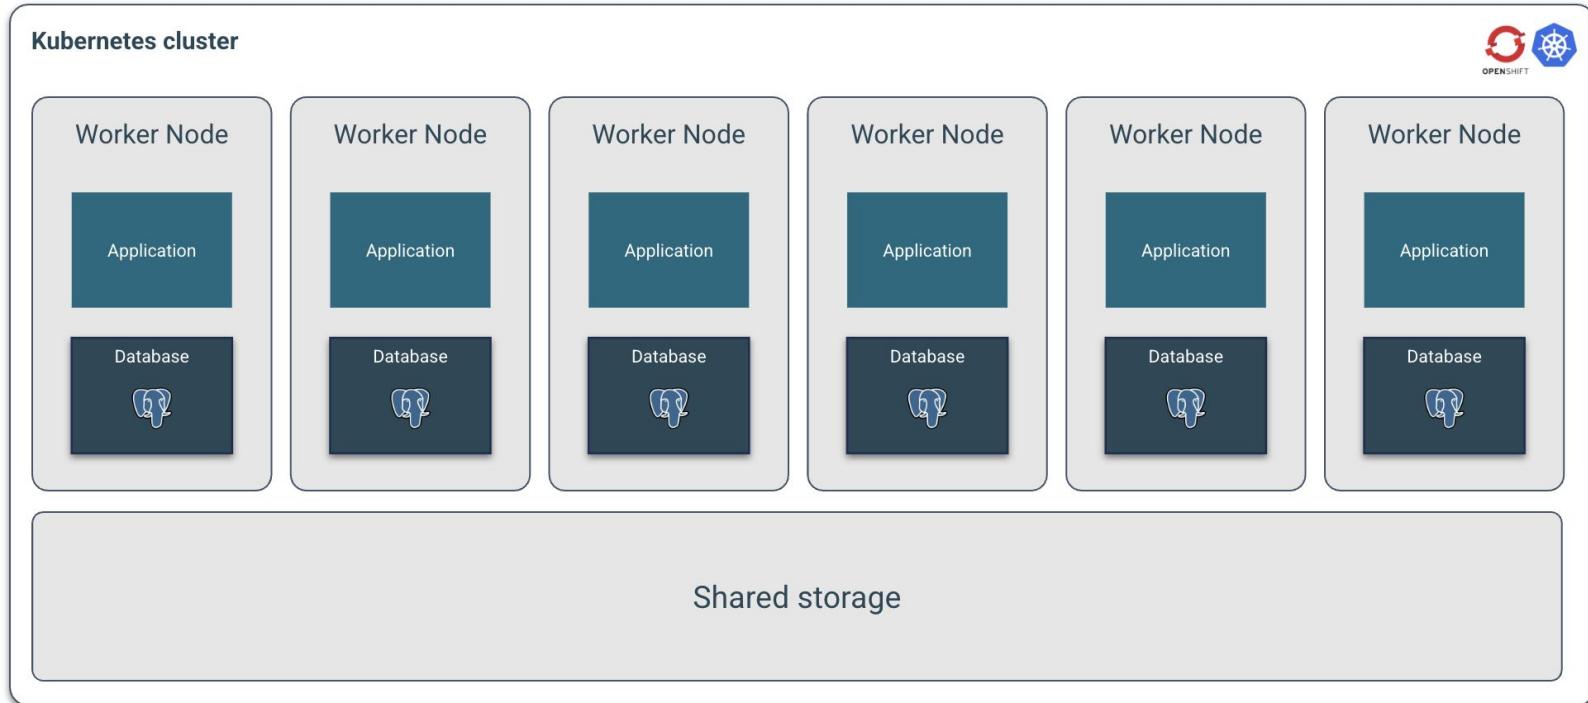


Two separate single data center Kubernetes clusters

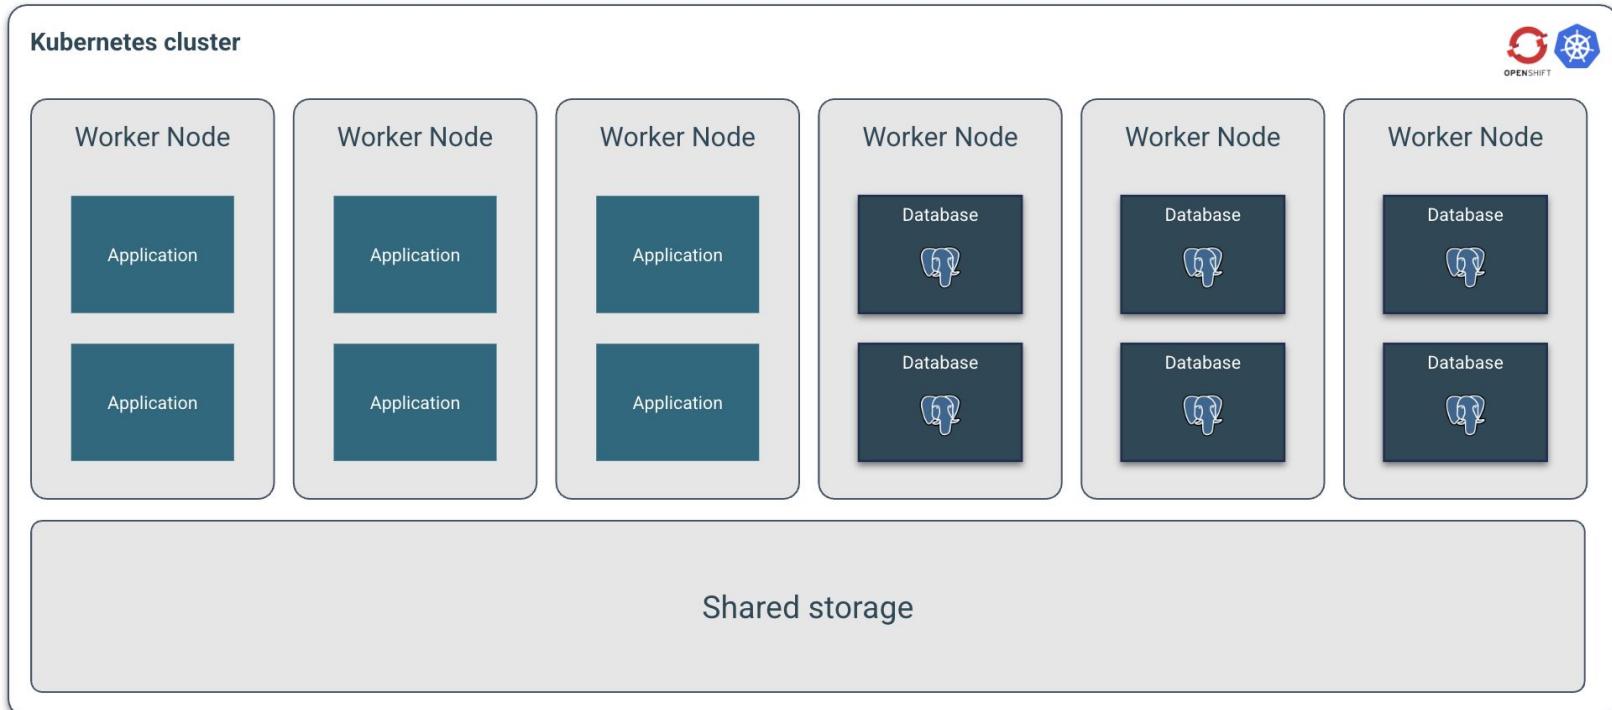
In case you cannot go beyond two data centers and you end up with two separate Kubernetes clusters, don't despair.



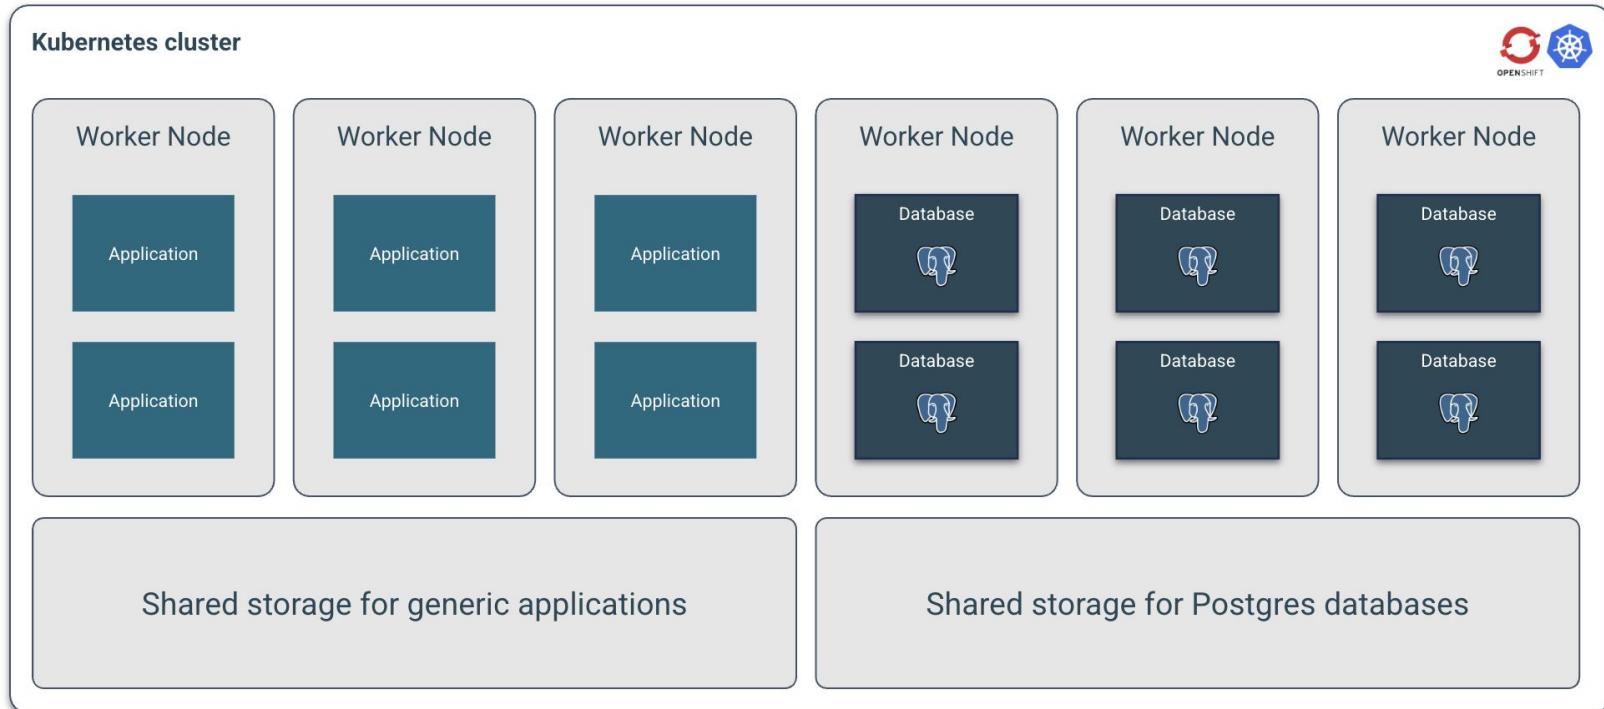
Shared workload, shared storage



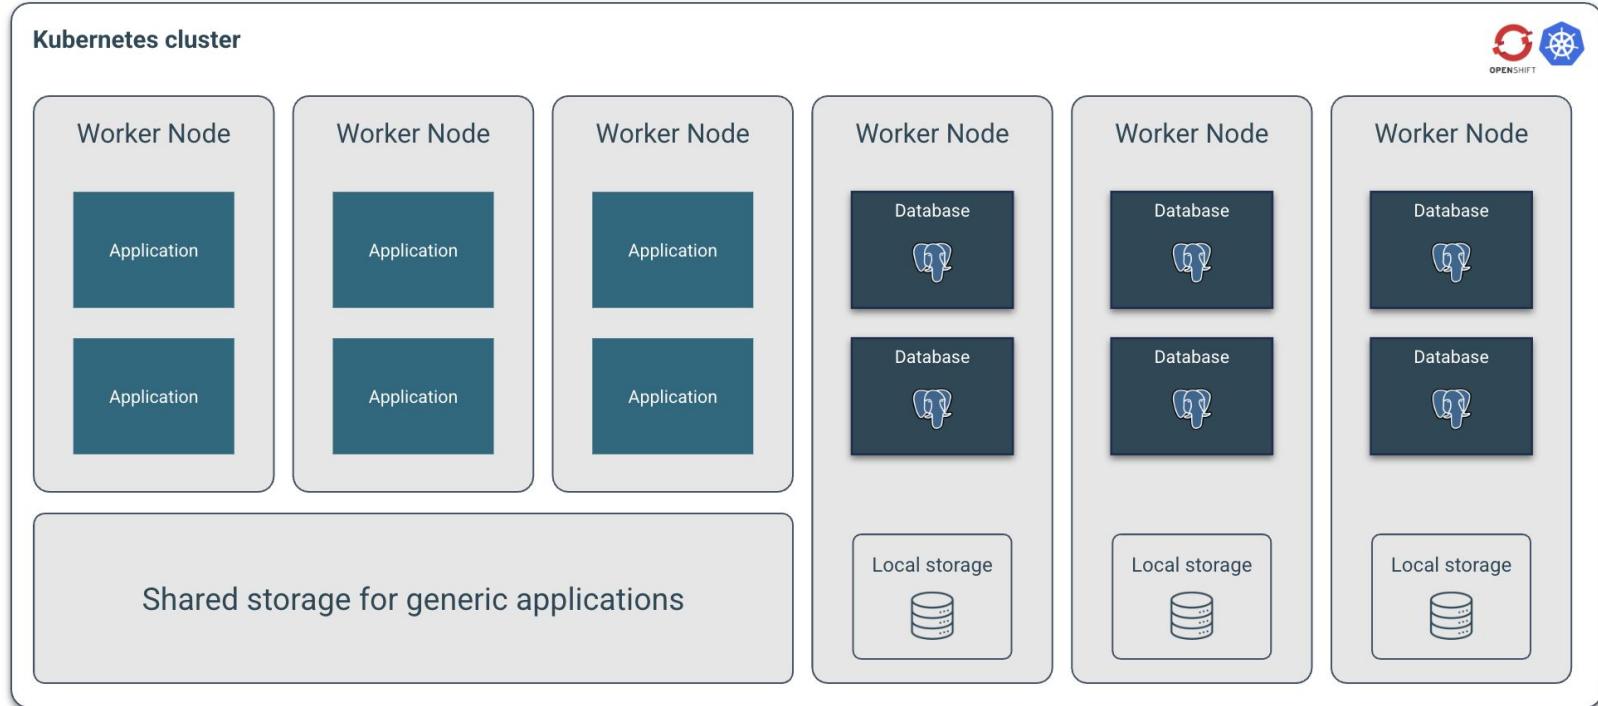
Shared workload, shared storage



Shared workload, shared storage



Shared workloads, local storage



Recommended architectures

<https://www.cncf.io/blog/2023/09/29/recommended-architectures-for-postgresql-in-kubernetes/>



Recommended architectures for PostgreSQL in Kubernetes

By Gabriele Bartolini

September 29, 2023

Member post by Gabriele Bartolini, VP of Cloud Native at EDB

"You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM", [tweeted Kelsey Hightower just a few months ago](#). Quite the opposite from what the former Google engineer and advocate said back in 2018 on Twitter: "Kubernetes supports stateful workloads; I don't."



Kelsey Hightower
@kelseyhightower

You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM. The biggest challenge is understanding that rubbing Kubernetes on Postgres won't turn it into Cloud SQL. [\[link\]](#)

Truth is that I agree with him now as much as I agreed with him back then. At that time, the holistic offering of storage capabilities in Kubernetes was still immature (local persistent volumes would become GA only the year after), the operator pattern – which in the meantime has proven to be crucial for stateful applications like databases – was yet to become widely accepted, and the [Data on Kubernetes Community](#) was more than two years away (second half of 2020).

Nowadays, the situation is completely different. And I am sure that many people who've worked hard in the last few years to bring stateful workloads in Kubernetes agree with me that Kelsey's recent powerful words will contribute to reversing the public perception and facilitate our mission – provided we keep doing great.



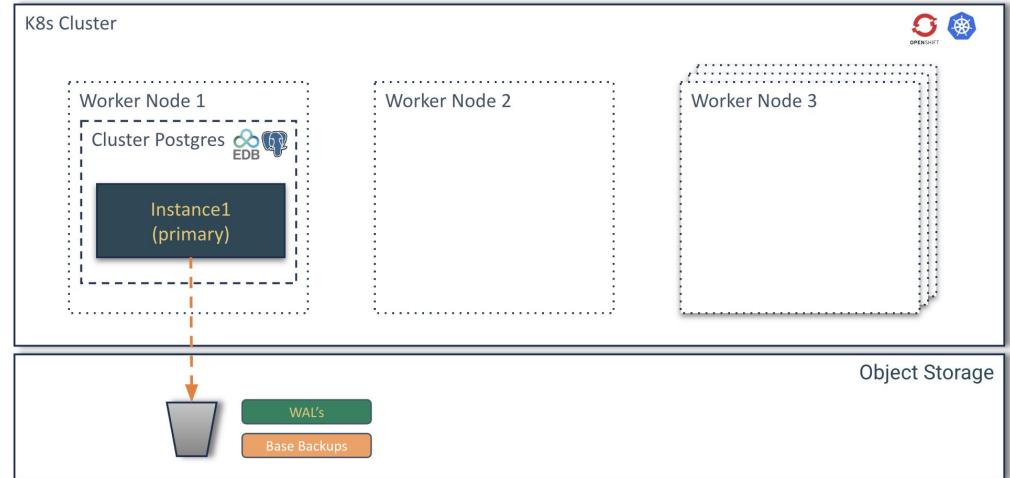
Use Cases



Use case 1 architecture

A single database is the simplest setup, involving one instance of a database server.

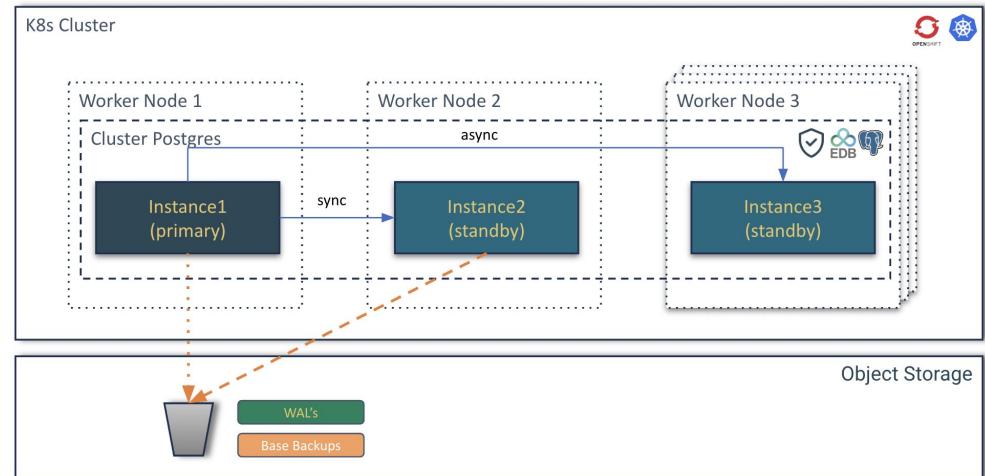
- Development and testing environments
- Small applications with low traffic
- Non-critical data analysis
- Applications with high tolerance for downtime
- Cost-sensitive projects



Use case 2 architecture

An HA database setup aims to minimize downtime by having redundant components. If one component fails, another takes over automatically or with minimal intervention. This usually involves techniques like clustering, replication, or mirroring within the same data center or availability zone.

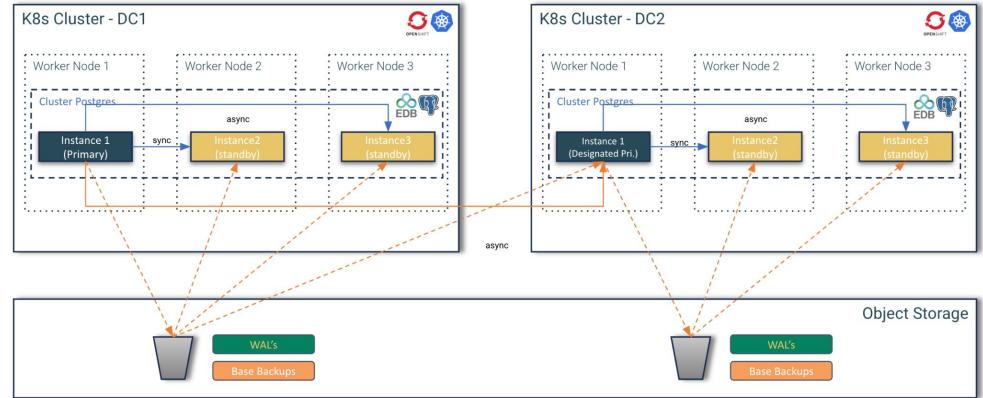
- Business critical Applications
- Applications with stringent SLAs
- Real-time systems
- Improving user experience
- Minimizing planned downtime



Use case 3 architecture

A DR database setup focuses on protecting data and ensuring business continuity in the event of a large-scale disaster affecting an entire data center or region (e.g., natural disasters, power outages, cyberattacks). This typically involves replicating data to a geographically separate location.

- Regulatory compliance
- Protecting against catastrophic data loss
- Ensuring business continuity for mission-critical systems



Use case 3 architecture

